

Section I - Setup

To set up, you will need two things: SQLiteStudio and the files for this class (which you likely have already if you are reading this document).

SQLiteStudio can be downloaded from its website: <https://sqlitestudio.pl/index.rvt?act=download>

The files for this class can be downloaded here: https://github.com/thomasniel/oreilly_advanced_sql_for_data

Section II - Subqueries, Unions, and Advanced Aggregations

2.1A - Scalar Subqueries

Get all orders on the last day there were orders

```
SELECT * FROM CUSTOMER_ORDER
WHERE ORDER_DATE = (SELECT MAX(ORDER_DATE) FROM CUSTOMER_ORDER)
```

2.1B - Aliasing Tables and Scalar Subquery Aggregation

Retrieving the average of quantity by each record's CUSTOMER_ID and PRODUCT_ID

```
SELECT CUSTOMER_ORDER_ID,
CUSTOMER_ID,
ORDER_DATE,
PRODUCT_ID,
QUANTITY,
(
    SELECT AVG(QUANTITY)
    FROM CUSTOMER_ORDER c2
    WHERE c2.CUSTOMER_ID = c1.CUSTOMER_ID
    AND c2.PRODUCT_ID = c1.PRODUCT_ID
) AS AVG_QUANTITY
FROM CUSTOMER_ORDER c1
```

Depending on how they are used, subqueries can be more expensive or less expensive than joins. Subqueries that generate a value for each record tend to be more expensive, like the example above.

2.2C - Multi-value Subqueries

Sometimes it can be helpful to leverage subqueries that return a set of values, rather than one scalar value. For instance, to query customer orders for customers in TX, we can save ourselves a join and use a subquery to get CUSTOMER_ID's that belong to customers in TX. Then we can leverage that with a WHERE and specify an IN condition:

```
SELECT CUSTOMER_ORDER_ID,  
CUSTOMER_ID,  
ORDER_DATE,  
PRODUCT_ID,  
QUANTITY  
  
FROM CUSTOMER_ORDER  
  
WHERE CUSTOMER_ID IN (  
    SELECT CUSTOMER_ID  
    FROM CUSTOMER  
    WHERE STATE = 'TX'  
)
```

2.3 - Derived Tables

A more efficient way to bring in averages by CUSTOMER_ID and PRODUCT_ID is by deriving a table of these averages, and joining to it.

```
SELECT CUSTOMER_ORDER_ID,  
CUSTOMER_ID,  
ORDER_DATE,  
PRODUCT_ID,  
QUANTITY,  
cust_avgs.avg_qty  
  
FROM CUSTOMER_ORDER  
INNER JOIN  
(  
    SELECT CUSTOMER_ID,  
    PRODUCT_ID,  
    AVG(QUANTITY) as avg_qty  
    FROM CUSTOMER_ORDER  
    GROUP BY 1, 2  
) cust_avgs  
ON CUSTOMER_ORDER.CUSTOMER_ID = cust_avgs.CUSTOMER_ID  
AND CUSTOMER_ORDER.PRODUCT_ID = cust_avgs.PRODUCT_ID
```

2.4 - Unions

To simply append two queries (with identical fields) together, put a `UNION ALL` between them.

```
SELECT
'FEB' AS MONTH,
PRODUCT.PRODUCT_ID,
PRODUCT_NAME,
SUM(PRICE * QUANTITY) AS REV
FROM PRODUCT LEFT JOIN CUSTOMER_ORDER
ON PRODUCT.PRODUCT_ID = CUSTOMER_ORDER.PRODUCT_ID

WHERE ORDER_DATE BETWEEN '2017-02-01' AND '2017-02-28'
GROUP BY 1,2,3

UNION ALL

SELECT
'MAR' AS MONTH,
PRODUCT.PRODUCT_ID,
PRODUCT_NAME,
SUM(PRICE * QUANTITY) AS REV
FROM PRODUCT LEFT JOIN CUSTOMER_ORDER
ON PRODUCT.PRODUCT_ID = CUSTOMER_ORDER.PRODUCT_ID

WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'
GROUP BY 1,2,3
```

Using `UNION` instead of `UNION ALL` will remove duplicates, which should not be necessary in this case.

You should strive not to use unions as they often encourage bad, inefficient SQL. Strive to use `CASE` statements or other tools instead. In this example, it would have been better to do this:

```
SELECT
CASE
    WHEN ORDER_DATE BETWEEN '2017-02-01' AND '2017-02-28' THEN 'FEB'
    WHEN ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31' THEN 'MAR'
END AS MONTH,
PRODUCT.PRODUCT_ID,
PRODUCT_NAME,
SUM(PRICE * QUANTITY) AS REV
FROM PRODUCT LEFT JOIN CUSTOMER_ORDER
ON PRODUCT.PRODUCT_ID = CUSTOMER_ORDER.PRODUCT_ID
```

```
WHERE ORDER_DATE BETWEEN '2017-02-01' AND '2017-03-31'  
GROUP BY 1,2,3
```

2.5 - GROUP CONCAT

A neat little trick you can do on some database platforms (like SQLite, MySQL, and PostgreSQL) is the `group_concat()` aggregate function. This will concatenate all values in a column as an aggregation, and can be used in conjunction with a GROUP BY like MIN, MAX, AVG, etc.

This shows a concatenated list of values of `PRODUCT_ID`'s ordered for each `ORDER_DATE`.

```
SELECT ORDER_DATE,  
group_concat(PRODUCT_ID) as product_ids_ordered  
  
FROM CUSTOMER_ORDER  
WHERE ORDER_DATE BETWEEN '2017-02-01' AND '2017-02-28'  
GROUP BY 1
```

Putting the `DISTINCT` keyword inside of it will only concatenate the `DISTINCT` product ID's.

```
SELECT ORDER_DATE,  
group_concat(DISTINCT PRODUCT_ID) as product_ids_ordered  
  
FROM CUSTOMER_ORDER  
WHERE ORDER_DATE BETWEEN '2017-02-01' AND '2017-02-28'  
GROUP BY 1
```

`GROUP_CONCAT` is a helpful function to compress the results into a single record, in a single cell, often in a reporting context.

Exercise 2-2

Bring in all records for `CUSTOMER_ORDER`, but also bring in the minimum and maximum quantities ever ordered each given `PRODUCT_ID` and `CUSTOMER_ID`.

ANSWER:

```
SELECT CUSTOMER_ORDER_ID,  
CUSTOMER_ID,  
ORDER_DATE,  
PRODUCT_ID,  
QUANTITY,  
min_qty,
```

```

max_qty

FROM CUSTOMER_ORDER
INNER JOIN
(
    SELECT CUSTOMER_ID,
    PRODUCT_ID,
    MIN(QUANTITY) AS min_qty,
    MAX(QUANTITY) AS max_qty
    FROM CUSTOMER_ORDER
    GROUP BY 1, 2
) min_and_max_ordered

ON CUSTOMER_ORDER.CUSTOMER_ID = min_and_max_ordered.CUSTOMER_ID
AND CUSTOMER_ORDER.PRODUCT_ID = min_and_max_ordered.PRODUCT_ID

```

Section III - Regular Expressions

Regular expressions are a powerful tool for qualifying complex text patterns. Their usage extends far outside of SQL and can be found on many technology platforms including Python, R, Java, .NET, LibreOffice, Alteryx, and Tableau.

While regular expressions can be used to split and search text, we will primarily be using it to match text, much like wildcards.

The REGEXP operator is used in SQLite for matching a text to a regex pattern. Like all boolean operations, it will return a 1 for true or 0 for false.

3.1 - Literals

Literals are characters in a regex pattern that have no special function, and represent that character verbatim. Numbers and letters are literals. For example, The regex TX will match the string TX

```
SELECT 'TX' REGEXP 'TX' --true
```

Some characters, as we have seen, have special functionality in a regex. If you want to use these characters as literals, sometimes you have to escape them with a preceding \. These characters include:

```
[^$.|?*( )
```

So to qualify a U.S. currency amount, you will need to escape the dollar sign \$ and the decimal place .

```
SELECT '$181.12' REGEXP '\$181\.12' -- true
```

3.1 - Qualifying Alphabetic and Numeric Ranges

A range is a valid set of values for a single character. For instance [A-Z] qualifies one uppercase alphabetic value for the range of letters A thru Z, so [A-Z] [A-Z] would qualify two uppercase alphabetic text values. This would match the text string TX.

```
SELECT 'TX' REGEXP '[A-Z] [A-Z]' --true
SELECT '45' REGEXP '[A-Z] [A-Z]' --false
SELECT 'T2' REGEXP '[A-Z] [0-3]' --true
SELECT 'T9' REGEXP '[A-Z0-9] [A-Z0-9]' --true
```

TX would not match [A-Z] [A-Z] [A-Z] though because it is not three characters.

```
SELECT 'TX' REGEXP '[A-Z] [A-Z] [A-Z]' --false
SELECT 'ASU' REGEXP '[A-Z] [A-Z] [A-Z]' --true
```

We can also specify certain characters, and they don't necessarily have to be ranges:

```
SELECT 'A6' REGEXP '[ATUX] [469]' --true
SELECT 'B8' REGEXP '[ATUX] [469]' --false
```

Conversely, we can negate a set of characters by starting the range with ^:

```
SELECT 'A6' REGEXP '[^ATUX] [^469]' --false
SELECT 'B8' REGEXP '[^ATUX] [^469]' --true
```

3.2 - Anchoring

If you don't want partial matches but rather full matches, you have to anchor the beginning and end of the String with ^ and \$ respectively.

For instance, [A-Z] [A-Z] would qualify with SMU. This is because it found two alphabetic characters within those three characters.

```
SELECT 'SMU' REGEXP '[A-Z] [A-Z]' --true
```

If you don't want that, you will need to qualify start and end anchors, which effectively demands a full match when both are used:

```
SELECT 'SMU' REGEXP '^ [A-Z] [A-Z] $' --false
```

You can also anchor to just the beginning or end of the string to check, for instance, if a string starts with a number followed by an alphabetic character:

```
SELECT '9FN' REGEXP '^ [0-9] [A-Z]' --true
SELECT 'RFX' REGEXP '^ [0-9] [A-Z]' --false
```

3.3 - Repeaters

Sometimes we simply want to qualify a repeated pattern in our regular expression. For example, this is redundant:

```
SELECT 'ASU' REGEXP '^[A-Z][A-Z][A-Z]$' --true
```

We can instead explicitly identify in curly brackets we want to repeat that alphabetic character 3 times, by following it with a {3}.

```
SELECT 'ASU' REGEXP '^[A-Z]{3}$' --true
```

We can also specify a min and max number of repetitions, such as a minimum of 2 but max of 3.

```
SELECT 'ASU' REGEXP '^[A-Z]{2,3}$' --true
```

```
SELECT 'TX' REGEXP '^[A-Z]{2,3}$' --true
```

Leaving the second argument blank will result in only requiring a minimum of repetitions:

```
SELECT 'A' REGEXP '^[A-Z]{2,}$' --false
```

```
SELECT 'ASDIKJFSKJJJXVJGTHEWIROQWERKJTX' REGEXP '^[A-Z]{2,}$' --true
```

To allow 1 or more repetitions, use the +. This will qualify with 1 or more alphanumeric characters.

```
SELECT 'ASDFJSKJ4892KSFJJ2843KJSNBKW' REGEXP '^[A-Z0-9]+$' --true
```

```
SELECT 'SDFJSDKJF/&SSDKJ$#SDFKSDFKJ' REGEXP '^[A-Z0-9]+$' --false
```

To allow 0 or more repetitions, use the *

```
SELECT 'ASDFJSKJ4892KSFJJ2843KJSNBKW' REGEXP '^[A-Z0-9]*' --true
```

```
SELECT '' REGEXP '^[A-Z0-9]*' --true
```

To allow 0 or 1 repetitions (an optional character), follow the item with a ?. This will allow two characters to be preceded with a number, but it doesn't have to:

```
SELECT '9FX' REGEXP '^[0-9]?[A-Z]{2}$' --true
```

```
SELECT 'AX' REGEXP '^[0-9]?[A-Z]{2}$' --true
```

You can use several repeaters for different clauses in a regex. Below, we qualify a string of alphabetic characters, a dash - followed by a string of numbers, and then another - with a string of alphabetic characters.

```
SELECT 'ASJSDFH-32423522-HUETHNB' REGEXP '^[A-Z]+-[0-9]+-[A-Z]+$' --true
```

3.4 Wildcards

A dot . represents any character, even whitespaces.

```
SELECT 'A-3' REGEXP '...' --true
```

You can also use it with repeaters to create broad wildcards for any number of characters.

```
SELECT 'A-3' REGEXP '.{3}' --true
SELECT 'A-3' REGEXP '.*' --true
SELECT 'A-3' REGEXP '.*' --true
```

`.*` is a common way to express qualifying any text.

3.5 Alternation and Grouping

You can group up parts of a regular expression using rounded parenthesis (), often to put a repeater on that entire group. For example, we can make the entire decimal part of a dollar amount optional:

```
'sql SELECT '181.12' REGEXP '^([0-9]+(\.[0-9]{2})?)?$' --true SELECT
'181' REGEXP '^([0-9]+(\.[0-9]{2})?)?$' --true
```

We can also qualify a string of letters, a slash /, and a string of numbers, but qualify any number of repetitions of this entire pattern:

```
SELECT 'WHISKY/23482374/ZULU/23423234/FOXTROT/6453' REGEXP '^([A-Z]+/[0-9]+/?)+$' --true
```

The pipe | operator functions as an alternator operator, or effectively an OR. It allows you to qualify any number of regular expressions where at least one of them must be true:

```
SELECT 'ALPHA' REGEXP '^(FOXTROT|ZULU|ALPHA|TANGO)$' --true
```

3.7 Using Regular Expressions in queries

Find all customers with a 3-4 digit street number. Note the literal space before the wildcard `.*`:

```
SELECT * FROM CUSTOMER
WHERE ADDRESS REGEXP '^[0-9]{3,4} .*$'
```

EXERCISE

Find all customers with an address ending in “Blvd” or “St”:

```
SELECT * FROM CUSTOMER
WHERE ADDRESS REGEXP '.*(Blvd|St)$'
```


Section IV - Advanced Joins

4.1 Inner Join Review

Using an INNER JOIN, you can view CUSTOMER and PRODUCT information with each CUSTOMER_ORDER.

```
SELECT CUSTOMER_ORDER_ID,  
CUSTOMER_NAME,  
ORDER_DATE,  
PRODUCT_ID,  
PRODUCT_NAME,  
PRODUCT_GROUP  
QUANTITY,  
PRICE  
  
FROM CUSTOMER_ORDER  
  
INNER JOIN CUSTOMER  
ON CUSTOMER_ORDER.CUSTOMER_ID = CUSTOMER.CUSTOMER_ID  
  
INNER JOIN PRODUCT  
ON CUSTOMER_ORDER.PRODUCT_ID = CUSTOMER_ORDER.PRODUCT_ID
```

4.2 Left Join Review

Total quantity sold of all products, null means no products were sold that day.

```
SELECT PRODUCT_ID,  
PRODUCT_NAME,  
SUM(QUANTITY) as total_quantity  
  
FROM PRODUCT LEFT JOIN CUSTOMER_ORDER  
  
ON PRODUCT.PRODUCT_ID = CUSTOMER_ORDER.PRODUCT_ID  
AND ORDER_DATE = '2017-03-01'  
  
GROUP BY 1, 2
```

4.3 Creating a Volatile Table

Herei show to create a volatile/temporary table of discount rules. This table will dispose at the end of each session. It is no different than a standard CREATE TABLE statement other than the TEMP keyword.

```

CREATE TEMP TABLE DISCOUNT (
    CUSTOMER_ID_REGEX VARCHAR (20) NOT NULL DEFAULT ('.*'),
    PRODUCT_ID_REGEX   VARCHAR (20) NOT NULL DEFAULT ('.*'),
    PRODUCT_GROUP_REGEX VARCHAR (30) NOT NULL DEFAULT ('.*'),
    STATE_REGEX        VARCHAR (30) NOT NULL DEFAULT ('.*'),
    DISCOUNT_RATE     DOUBLE      NOT NULL
);

INSERT INTO DISCOUNT (STATE_REGEX, DISCOUNT_RATE) VALUES ('LA|OK', 0.20);
INSERT INTO DISCOUNT (PRODUCT_GROUP_REGEX, STATE_REGEX, DISCOUNT_RATE) VALUES ('BETA|GAMMA', 0.10);
INSERT INTO DISCOUNT (PRODUCT_ID_REGEX, CUSTOMER_ID_REGEX, DISCOUNT_RATE) VALUES ('^[379]$', 0.15);

```

Note you can also create a temporary (or permanent) table from a SELECT query. This is helpful to persist expensive query results and reuse it multiple times during a session. SQLite is a bit more convoluted to do this than other platforms:

```

CREATE TEMP TABLE ORDER_TOTALS_BY_DATE AS
WITH ORDER_TOTALS_BY_DATE AS (
    SELECT ORDER_DATE,
           SUM(QUANTITY) AS TOTAL_QUANTITY
    FROM CUSTOMER_ORDER
    GROUP BY 1
)
SELECT * FROM ORDER_TOTALS_BY_DATE

```

4.4 Joining with Regular Expressions

Left-joining to the temporary table and qualifying on the regular expressions for each respective field allows us to apply the discounts to each CUSTOMER_ORDER as specified.

```

SELECT CUSTOMER_ORDER.*,
       DISCOUNT_RATE,
       PRICE * (1 - DISCOUNT_RATE) AS DISCOUNTED_PRICE

FROM CUSTOMER_ORDER
INNER JOIN CUSTOMER
ON CUSTOMER_ORDER.CUSTOMER_ID = CUSTOMER.CUSTOMER_ID

INNER JOIN PRODUCT
ON CUSTOMER_ORDER.PRODUCT_ID = PRODUCT.PRODUCT_ID

LEFT JOIN DISCOUNT
ON CUSTOMER_ORDER.CUSTOMER_ID REGEXP DISCOUNT.CUSTOMER_ID_REGEX
AND CUSTOMER_ORDER.PRODUCT_ID REGEXP DISCOUNT.PRODUCT_ID_REGEX

```

```

AND PRODUCT.PRODUCT_GROUP REGEXP DISCOUNT.PRODUCT_GROUP_REGEX
AND CUSTOMER.STATE REGEXP DISCOUNT.STATE_REGEX

```

```

WHERE ORDER_DATE BETWEEN '2017-03-26' AND '2017-03-31'

```

If you expect records to possibly get multiple discounts, then sum the discounts and GROUP BY everything else:

```

SELECT CUSTOMER_ORDER_ID,
CUSTOMER_NAME,
STATE,
ORDER_DATE,
PRODUCT_ID,
PRODUCT_NAME,
PRODUCT_GROUP,
QUANTITY,
PRICE,
SUM(DISCOUNT_RATE) as TOTAL_DISCOUNT_RATE,
PRICE * (1 - SUM(DISCOUNT_RATE)) AS DISCOUNTED_PRICE

FROM CUSTOMER_ORDER
INNER JOIN CUSTOMER
ON CUSTOMER_ORDER.CUSTOMER_ID = CUSTOMER.CUSTOMER_ID

INNER JOIN PRODUCT
ON CUSTOMER_ORDER.PRODUCT_ID = CUSTOMER_ORDER.PRODUCT_ID

LEFT JOIN DISCOUNT
ON CUSTOMER_ORDER.CUSTOMER_ID REGEXP DISCOUNT.CUSTOMER_ID_REGEX
AND CUSTOMER_ORDER.PRODUCT_ID REGEXP DISCOUNT.PRODUCT_ID_REGEX
AND PRODUCT.PRODUCT_GROUP REGEXP DISCOUNT.PRODUCT_GROUP_REGEX
AND CUSTOMER.STATE REGEXP DISCOUNT.STATE_REGEX

WHERE ORDER_DATE BETWEEN '2017-03-26' AND '2017-03-31'

GROUP BY 1,2,3,4,5,6,7,8

```

4.5 Self Joins

We can join a table to itself by invoking it twice with two aliases. This can be useful, for example, to look up the previous day's order quantity (if any) for a given CUSTOMER_ID and PRODUCT_ID:

```

SELECT o1.CUSTOMER_ORDER_ID,

```

```

o1.CUSTOMER_ID,
o1.PRODUCT_ID,
o1.ORDER_DATE,
o1.QUANTITY,
o2.QUANTITY AS PREV_DAY_QUANTITY

FROM CUSTOMER_ORDER o1
LEFT JOIN CUSTOMER_ORDER o2

ON o1.CUSTOMER_ID = o2.CUSTOMER_ID
AND o1.PRODUCT_ID = o2.PRODUCT_ID
AND o2.ORDER_DATE = date(o1.ORDER_DATE, '-1 day')

WHERE o1.ORDER_DATE BETWEEN '2017-03-05' AND '2017-03-11'

```

Note if you want to get the previous quantity ordered for that record's given CUSTOMER_ID and PRODUCT_ID, even if it wasn't strictly the day before, you can use a subquery instead that qualifies previous dates and orders them descending. Then you can use LIMIT 1 to grab the most recent at the top.

```

SELECT ORDER_DATE,
PRODUCT_ID,
CUSTOMER_ID,
QUANTITY,
(
    SELECT QUANTITY
    FROM CUSTOMER_ORDER c2
    WHERE c1.ORDER_DATE > c2.ORDER_DATE
    AND c1.PRODUCT_ID = c2.PRODUCT_ID
    AND c1.CUSTOMER_ID = c2.CUSTOMER_ID
    ORDER BY ORDER_DATE DESC
    LIMIT 1
) as PREV_QTY
FROM CUSTOMER_ORDER c1

```

4.6 Cross Joins

Sometimes it can be helpful to generate a “cartesian product”, or every possible combination between two or more data sets using a CROSS JOIN. This is often done to generate a data set that fills in gaps for another query. Not every calendar date has orders, nor does every order date have an entry for every product, as shown in this query:

```

SELECT ORDER_DATE,
PRODUCT_ID,
SUM(QUANTITY) as TOTAL_QTY

```

```
FROM CUSTOMER_ORDER
```

```
GROUP BY 1, 2
```

We should use a cross join to resolve this problem. For instance, we can leverage a `CROSS JOIN` query to generate every possible combination of `PRODUCT_ID` and `CUSTOMER_ID`.

```
SELECT
CUSTOMER_ID,
PRODUCT_ID
FROM CUSTOMER
CROSS JOIN PRODUCT
```

In this case we should bring in `CALENDAR_DATE` and cross join it with `PRODUCT_ID` to get every possible combination of calendar date and product. Note the `CALENDAR_DATE` comes from the `CALENDAR` table, which acts as a simple list of consecutive calendar dates. We should only filter the calendar to a date range of interest, like 2017-01-01 and 2017-03-31.

```
SELECT
CALENDAR_DATE,
PRODUCT_ID
FROM PRODUCT
CROSS JOIN CALENDAR
WHERE CALENDAR_DATE BETWEEN '2017-01-01' and '2017-03-31'
```

Then we can `LEFT JOIN` to our previous query to get every product quantity sold by calendar date, even if there were no orders that day:

```
SELECT CALENDAR_DATE,
all_combos.PRODUCT_ID,
TOTAL_QTY

FROM
(
    SELECT
    CALENDAR_DATE,
    PRODUCT_ID
    FROM PRODUCT
    CROSS JOIN CALENDAR
    WHERE CALENDAR_DATE BETWEEN '2017-01-01' and '2017-03-31'
) all_combos

LEFT JOIN
(
    SELECT ORDER_DATE,
    PRODUCT_ID,
```

```

SUM(QUANTITY) as TOTAL_QTY

FROM CUSTOMER_ORDER

GROUP BY 1, 2
) totals

ON all_combos.CALENDAR_DATE = totals.ORDER_DATE
AND all_combos.PRODUCT_ID = totals.PRODUCT_ID

```

4.7 Comparative Joins

Note also you can use comparison operators in joins. For instance, we can self-join to create rolling quantity totals and generate a cartesian product on previous dates to the current order, and then sum those quantities. It is much easier to use windowing functions for this purpose though, which is covered in the next section.

```

SELECT c1.ORDER_DATE,
c1.PRODUCT_ID,
c1.CUSTOMER_ID,
c1.QUANTITY,
SUM(c2.QUANTITY) as ROLLING_QTY

FROM CUSTOMER_ORDER c1 INNER JOIN CUSTOMER_ORDER c2
ON c1.PRODUCT_ID = c2.PRODUCT_ID
AND c1.CUSTOMER_ID = c2.CUSTOMER_ID
AND c1.ORDER_DATE >= c2.ORDER_DATE

GROUP BY 1, 2, 3, 4

```

Exercise 4

For every CALENDAR_DATE and CUSTOMER_ID, show the total QUANTITY ordered for the date range of 2017-01-01 to 2017-03-31:

ANSWER:

```

SELECT CALENDAR_DATE,
all_combos.CUSTOMER_ID,
TOTAL_QTY

FROM
(
SELECT

```

```

    CALENDAR_DATE,
    CUSTOMER_ID
  FROM CUSTOMER
  CROSS JOIN CALENDAR
  WHERE CALENDAR_DATE BETWEEN '2017-01-01' and '2017-03-31'
) all_combos

LEFT JOIN
(
  SELECT ORDER_DATE,
  CUSTOMER_ID,
  SUM(QUANTITY) as TOTAL_QTY

  FROM CUSTOMER_ORDER

  GROUP BY 1, 2
) totals

ON all_combos.CALENDAR_DATE = totals.ORDER_DATE
AND all_combos.CUSTOMER_ID = totals.CUSTOMER_ID

```

Section V - Windowing

Windowing functions allow you to create contextual aggregations in ways much more flexible than GROUP BY. Many major database platforms support windowing functions, including:

- Oracle
- Teradata
- PostgreSQL
- SQL Server
- Apache Spark SQL

These platforms notably do not have windowing functions:

- MySQL
- SQLite
- MariaDB

Since SQLite does not support windowing functions, we are going to use PostgreSQL. While PostgreSQL is free and open-source, there are a few steps in getting it set up. Therefore to save time we are going to use Rextester, a web-based client that can run PostgreSQL queries.

http://rextester.com/1/postgresql_online_compiler

In the resources for this class, you should find a “customer_order.sql” file which

can be opened with any text editor. Inside you will see some SQL commands to create and populate a `CUSTOMER_ORDER` table and then `SELECT` all the records from it. Copy/Paste the contents to Rextester and then click the “Run it (F8)” button.

Notice it will create the table and populate it, and the final `SELECT` query will execute and display the results. Note that the table is not persisted after the operation finishes, so you will need to precede each `SELECT` exercise with this table creation and population before your `SELECT`.

5.1 PARTITION BY

Sometimes it can be helpful to create a contextual aggregation for each record in a query. Windowing functions can make this much easier and save us a lot of subquery work.

For instance, it may be helpful to not only get each `CUSTOMER_ORDER` for the month of MARCH, but also the maximum quantity that customer purchased for that `PRODUCT_ID`. We can do that with an `OVER PARTITION BY` combined with the `MAX()` function.

```
SELECT CUSTOMER_ORDER_ID,  
CUSTOMER_ID,  
ORDER_DATE,  
PRODUCT_ID,  
QUANTITY,  
MAX(QUANTITY) OVER(PARTITION BY PRODUCT_ID, CUSTOMER_ID) as MAX_PRODUCT_QTY_ORDERED  
  
FROM CUSTOMER_ORDER  
  
WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'  
  
ORDER BY CUSTOMER_ORDER_ID
```

Each `MAX_PRODUCT_QTY_ORDERED` will only be the maximum `QUANTITY` of that given record's `PRODUCT_ID` and `CUSTOMER_ID`. The `WHERE` will also filter that scope to only within MARCH.

You can have multiple windowed fields in a query. Below, we get a `MIN`, `MAX`, and `AVG` for that given window.

```
SELECT CUSTOMER_ORDER_ID,  
CUSTOMER_ID,  
ORDER_DATE,  
PRODUCT_ID,  
QUANTITY,  
MIN(QUANTITY) OVER(PARTITION BY PRODUCT_ID, CUSTOMER_ID) as MIN_PRODUCT_QTY_ORDERED,
```



```

MAX(QUANTITY) OVER(PARTITION BY PRODUCT_ID, CUSTOMER_ID) as MAX_PRODUCT_QTY_ORDERED,
AVG(QUANTITY) OVER(PARTITION BY PRODUCT_ID, CUSTOMER_ID) as AVG_PRODUCT_QTY_ORDERED

FROM CUSTOMER_ORDER

WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'

ORDER BY CUSTOMER_ORDER_ID

```

When you are declaring your window redundantly, you can reuse it using a WINDOW declaration, which goes between the WHERE and the ORDER BY.

```

SELECT CUSTOMER_ORDER_ID,
CUSTOMER_ID,
ORDER_DATE,
PRODUCT_ID,
QUANTITY,
MIN(QUANTITY) OVER(w) as MIN_PRODUCT_QTY_ORDERED,
MAX(QUANTITY) OVER(w) as MAX_PRODUCT_QTY_ORDERED,
AVG(QUANTITY) OVER(w) as AVG_PRODUCT_QTY_ORDERED

FROM CUSTOMER_ORDER

WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'

WINDOW w AS (PARTITION BY PRODUCT_ID, CUSTOMER_ID)

ORDER BY CUSTOMER_ORDER_ID

```

5.2 ORDER BY

You can also use an ORDER BY in your window to only consider values that comparatively come before that record.

5.2A USING ORDER BY

For instance, you can get a ROLLING_TOTAL of the QUANTITY by ordering by the ORDER_DATE.

```

SELECT CUSTOMER_ORDER_ID,
CUSTOMER_ID,
ORDER_DATE,
PRODUCT_ID,
QUANTITY,
SUM(QUANTITY) OVER (ORDER BY ORDER_DATE) AS ROLLING_QUANTITY

```

```

FROM CUSTOMER_ORDER

WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'

ORDER BY CUSTOMER_ORDER_ID

```

Note you can precede the `ORDER BY` clause with a `DESC` keyword to window in the opposite direction.

5.2B Ordering and Bounds

Above, notice our example output has the same rolling total for all records on a given date. This is because the `ORDER BY` in a window function by default does a logical boundary, which in this case is the `ORDER_DATE`. This means it is rolling up everything on that `ORDER_DATE` and previous to it. A side effect is all records with the same `ORDER_DATE` are going to get the same rolling total.

This is the default behavior our query did previously:

```

SELECT CUSTOMER_ORDER_ID,
CUSTOMER_ID,
ORDER_DATE,
PRODUCT_ID,
QUANTITY,
SUM(QUANTITY) OVER (ORDER BY ORDER_DATE RANGE BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) AS
FROM CUSTOMER_ORDER

WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'

ORDER BY CUSTOMER_ORDER_ID

```

If you want to incrementally roll the quantity by each row's physical order (not logical order by the entire `ORDER_DATE`), you can use `ROWS BETWEEN` instead of `RANGE BETWEEN`.

```

SELECT CUSTOMER_ORDER_ID,
CUSTOMER_ID,
ORDER_DATE,
PRODUCT_ID,
QUANTITY,
SUM(QUANTITY) OVER (ORDER BY ORDER_DATE ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW) AS
FROM CUSTOMER_ORDER

WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'

```

ORDER BY CUSTOMER_ORDER_ID

Note the **AND CURRENT ROW** is a default, so you can shorthand it like this:

SUM(QUANTITY) OVER (ORDER BY ORDER_DATE ROWS UNBOUNDED PRECEDING) AS ROLLING_QUANTITY

In this particular example, you could have avoided using a physical boundary by specifying your window with an **ORDER BY CUSTOMER_ORDER_ID**. But we covered the previous strategy anyway to see how to execute physical boundaries. Here is an excellent overview of windowing functions and bounds: <http://mysqlserverteam.com/mysql-8-0-2-introducing-window-functions/>

5.3 Mixing PARTITION BY / ORDER BY

We can combine the **PARTITION BY / ORDER BY** to create rolling aggregations partitioned on certain fields.

5.3A Simple MAX over PARTITION and ORDER BY

For example, for each record we can get the max quantity ordered up to that date

```
SELECT CUSTOMER_ORDER_ID,  
CUSTOMER_ID,  
ORDER_DATE,  
PRODUCT_ID,  
QUANTITY,  
MAX(QUANTITY) OVER(PARTITION BY PRODUCT_ID, CUSTOMER_ID ORDER BY ORDER_DATE) as MAX_TO_DATE,  
  
FROM CUSTOMER_ORDER  
  
WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'  
  
ORDER BY CUSTOMER_ORDER_ID
```

5.3B Rolling Total Quantity by PRODUCT_ID, with physical boundary

```
SELECT CUSTOMER_ORDER_ID,  
ORDER_DATE,  
CUSTOMER_ID,  
PRODUCT_ID,  
QUANTITY,
```

```
SUM(QUANTITY) OVER(PARTITION BY PRODUCT_ID ORDER BY ORDER_DATE ROWS UNBOUNDED PRECEDING) as
```

```
FROM CUSTOMER_ORDER
WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'
```

You need to be very careful mixing PARTITION BY with an ORDER BY that uses a physical boundary! If you sort the results, it can get confusing very quickly because you lose that physical ordered context.

```
SELECT CUSTOMER_ORDER_ID,
ORDER_DATE,
CUSTOMER_ID,
PRODUCT_ID,
QUANTITY,
SUM(QUANTITY) OVER(PARTITION BY PRODUCT_ID ORDER BY ORDER_DATE) as total_qty_for_customer_ar
```

```
FROM CUSTOMER_ORDER
WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'
```

```
ORDER BY ORDER_DATE
```

5.4 Rolling Windows

You can also use movable windows to create moving aggregations. For instance, you can create a six-day rolling average (3 days before, 3 days after).

Note that PostgreSQL does not support this but MySQL, Teradata, and a few other platforms do.

```
SELECT CUSTOMER_ORDER_ID,
CUSTOMER_ID,
ORDER_DATE,
PRODUCT_ID,
QUANTITY,
AVG(QUANTITY) OVER (ORDER BY ORDER_DATE RANGE BETWEEN 3 PRECEDING AND 3 FOLLOWING) AS SIX_D
```

```
FROM CUSTOMER_ORDER
```

```
WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'
```

```
ORDER BY CUSTOMER_ORDER_ID
```

EXERCISE

For the month of March, bring in the rolling sum of quantity ordered (up to each ORDER_DATE) by CUSTOMER_ID and PRODUCT_ID.

```
SELECT CUSTOMER_ORDER_ID,  
ORDER_DATE,  
CUSTOMER_ID,  
PRODUCT_ID,  
QUANTITY,  
SUM(QUANTITY) OVER(PARTITION BY CUSTOMER_ID, PRODUCT_ID) as total_qty_for_customer_and_produ  
  
FROM CUSTOMER_ORDER  
WHERE ORDER_DATE BETWEEN '2017-03-01' AND '2017-03-31'  
  
ORDER BY CUSTOMER_ORDER_ID
```

Section VI - SQL with Python, R, and Java

6.1A Using SQL with Python

When doing SQL with Python, you want to use SQLAlchemy. Below, we query and loop through the CUSTOMER table which is returned as an iteration of Tuples:

```
from sqlalchemy import create_engine, text  
  
engine = create_engine('sqlite:///thunderbird_manufacturing.db')  
conn = engine.connect()  
  
stmt = text("SELECT * FROM CUSTOMER")  
results = conn.execute(stmt)  
  
for r in results:  
    print(r)
```

6.1B Using SQL with Python

You can package up interactions with a database into helper functions. Below, we create a function called `get_all_customers()` which returns the results as a List of tuples:

```
from sqlalchemy import create_engine, text  
  
engine = create_engine('sqlite:///thunderbird_manufacturing.db')
```

```

conn = engine.connect()

def get_all_customers():
    stmt = text("SELECT * FROM CUSTOMER")
    return list(conn.execute(stmt))

print(get_all_customers())

```

6-1C Using SQL with Python

If you want to pass parameters to a query, mind to not insert parameters directly so you don't accidentally introduce SQL injection. Below, we create a helper function that retrieves a customer for a given ID from a database.

```

from sqlalchemy import create_engine, text

engine = create_engine('sqlite:///thunderbird_manufacturing.db')
conn = engine.connect()

def get_all_customers():
    stmt = text("SELECT * FROM CUSTOMER")
    return list(conn.execute(stmt))

def customer_for_id(customer_id):
    stmt = text("SELECT * FROM CUSTOMER WHERE CUSTOMER_ID = :id")
    return conn.execute(stmt, id=customer_id).first()

print(customer_for_id(3))

```

You can also use these functions to update data.

6.2 Using SQL with R

Here is how to run a SQL query in R, and save the results to a matrix.

```

setwd('c:\\my_folder')

library(DBI)
library(RSQLite)

db <- dbConnect(SQLite(), dbname='thunderbird_manufacturing.db')

```

```
myQuery <- dbSendQuery(db, "SELECT * FROM CUSTOMER")

my_data <- dbFetch(myQuery, n = -1)

dbClearResult(myQuery)

print(my_data)

remove(myQuery)
dbDisconnect(db)
```

You can get detailed information on how to work with R and SQL in the official DBI documentation: * DBI Interface: <https://cran.r-project.org/web/packages/DBI/index.html> * DBI PDF: <https://cran.r-project.org/web/packages/DBI/DBI.pdf>

6-3 Using SQL with Java/Scala/Kotlin

There are many solutions to make a Java, Scala, or Kotlin application work with a Java database. The vanilla way we will learn is to use JDBC (Java Database Connection).

Keep in mind there are many solutions and libraries that abstract away SQL operations, which can be good or bad depending on how much control you want to maintain:

- **Hibernate** - ORM technology that's been around since 2001 and has a mature implementation. However, Hibernate is notorious for its strange loading mechanisms, and can be a hindrance if you want to maintain control of how and when data is loaded from a database.
- **jOOQ** - A more modern (but commercial) ORM that fluently allows working with databases in a type-safe manner.
- **Speedment** - Another fast turnaround, fluent API that compiles pure Java code from table schemas to work with databases.

If you are going to go the vanilla JDBC route, it is a good idea to use a connection pool so you can persist and reuse several connections safely in a multithreaded environment. HikariCP is a leading option to achieve this and provides an optimal **DataSource** implementation, which is Java's recommended interface for a database connection pool.

A helpful resource to learning how to work with JDBC is Jenkov's in-depth tutorial: <http://tutorials.jenkov.com/jdbc/index.html>

6.3A - Selecting Data with JDBC and HikariCP

To connect to a database using JDBC and HikariCP, you will need the appropriate JDBC drivers for your database platform (e.g. SQLite) as well as Hikari-CP.

```
dependencies {  
    compile 'org.xerial:sqlite-jdbc:3.19.3'  
    compile 'com.zaxxer:HikariCP:2.6.3'  
    compile 'org.slf4j:slf4j-simple:1.7.25'  
}
```

Below, we create a simple Java application that creates a Hikari data source with a minimum of 1 connection and a maximum of 5. Then we create a query and loop through it's ResultSet.

```
import com.zaxxer.hikari.HikariConfig;  
import com.zaxxer.hikari.HikariDataSource;  
  
import java.sql.Connection;  
import java.sql.ResultSet;  
import java.sql.Statement;  
  
public class Launcher {  
  
    public static void main(String[] args) {  
  
        try {  
            HikariConfig config = new HikariConfig();  
            config.setJdbcUrl("jdbc:sqlite:/c:/git/oreilly_advanced_sql_for_data/thunderbird.sqlite");  
            config.setMinimumIdle(1);  
            config.setMaximumPoolSize(5);  
  
            HikariDataSource ds = new HikariDataSource(config);  
  
            Connection conn = ds.getConnection();  
            Statement stmt = conn.createStatement();  
            ResultSet rs = stmt.executeQuery("SELECT * from CUSTOMER");  
  
            while (rs.next()) {  
                System.out.println(rs.getInt("CUSTOMER_ID") + " " + rs.getString("CUSTOMER_NAME"));  
            }  
  
            //release connection back to pool  
            conn.close();  
  
        } catch (Exception e) {  
            e.printStackTrace();  
        }  
    }  
}
```



```

    }
}
}

```

6.3B Passing parameters

If you need to pass parameters to your SQL query, avoid concatenating the values into the SQL string otherwise you will put your application at risk for SQL injection. Instead, use a `PreparedStatement` to safely inject the parameters:

```

import com.zaxxer.hikari.HikariConfig;
import com.zaxxer.hikari.HikariDataSource;

import java.sql.Connection;
import java.sql.PreparedStatement;
import java.sql.ResultSet;

public class Launcher {

    public static void main(String[] args) {

        try {
            HikariConfig config = new HikariConfig();
            config.setJdbcUrl("jdbc:sqlite:/c:/git/oreilly_advanced_sql_for_data/thunderbird.sqlite");
            config.setMinimumIdle(1);
            config.setMaximumPoolSize(5);

            HikariDataSource ds = new HikariDataSource(config);

            Connection conn = ds.getConnection();

            // Create a PreparedStatement and populate parameter
            PreparedStatement stmt = conn.prepareStatement("SELECT * from CUSTOMER WHERE CUSTOMER_ID = ?");
            stmt.setInt(1,3);

            ResultSet rs = stmt.executeQuery();

            while (rs.next()) {
                System.out.println(rs.getInt("CUSTOMER_ID") + " " + rs.getString("CUSTOMER_NAME"));
            }

            //release connection back to pool
            conn.close();

        } catch (Exception e) {

```

```

        e.printStackTrace();
    }
}

```

6-3C Writing Data

You can also use a `PreparedStatement` to execute updates against the database. The `PreparedStatement` has more advanced features like batching to write large volumes of data, but here is how to insert a single record.

```

import com.zaxxer.hikari.HikariConfig;
import com.zaxxer.hikari.HikariDataSource;

import java.math.BigDecimal;
import java.sql.Connection;
import java.sql.PreparedStatement;

public class Launcher {

    public static void main(String[] args) {

        try {
            HikariConfig config = new HikariConfig();
            config.setJdbcUrl("jdbc:sqlite:/c:/git/oreilly_advanced_sql_for_data/thunderbird.db");
            config.setMinimumIdle(1);
            config.setMaximumPoolSize(5);

            HikariDataSource ds = new HikariDataSource(config);

            Connection conn = ds.getConnection();

            // Create a PreparedStatement and populate parameter
            PreparedStatement stmt =
                conn.prepareStatement("INSERT INTO PRODUCT (PRODUCT_NAME,PRODUCT_GROUP,PRICE) VALUES (?, ?, ?)");

            stmt.setString(1,"Kry Kall");
            stmt.setString(2,"BETA");
            stmt.setBigDecimal(3, BigDecimal.valueOf(35.0));

            stmt.executeUpdate();

            //release connection back to pool
            conn.close();
        } catch (Exception e) {
            e.printStackTrace();
        }
    }
}

```

```
    } catch (Exception e) {  
        e.printStackTrace();  
    }  
}
```